

About Us



Josh O'Rourke



Kevin Kingsbury



Scott Walters

Agenda

- Data Architecture
- Airflow Overview
- Airflow Concepts
- Airflow UI
- Demo
- Q&A

Data Architecture

A Naive Approach

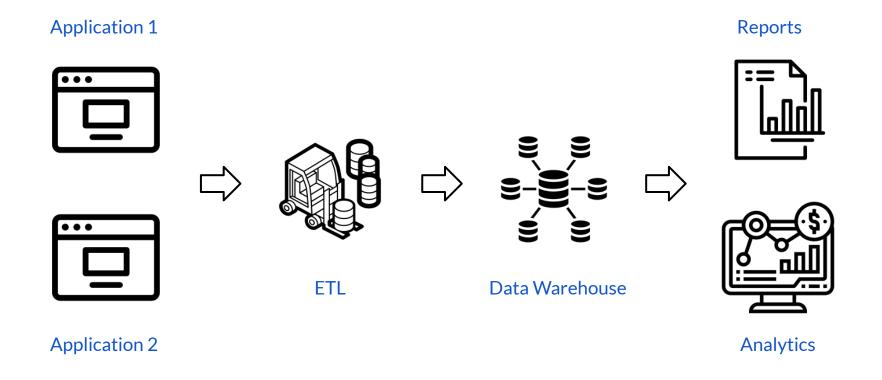
In the beginning there was Cron.

We had one job. It ran at 1 am.

And it was good.

Pete Owlett, PyData London 2016 Lessons from 6 Months of Using Luigi in Production

Data Architecture



Problems

- **Failures:** How do we retry when failure happens?
- Monitoring: Are we notified about errors? How long does the job run?
- **Dependencies**: What happens if upstream data is missing?
- Scalability: How do we coordinate cron jobs on different machines?
- **Deployment**: How often do we need to deploy changes? Is it easy?
- **Historic Data**: Can we backfill historical data?
- Logging: Do we have unified logs?

Modern Data Architecture

Modern Architecture







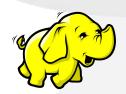








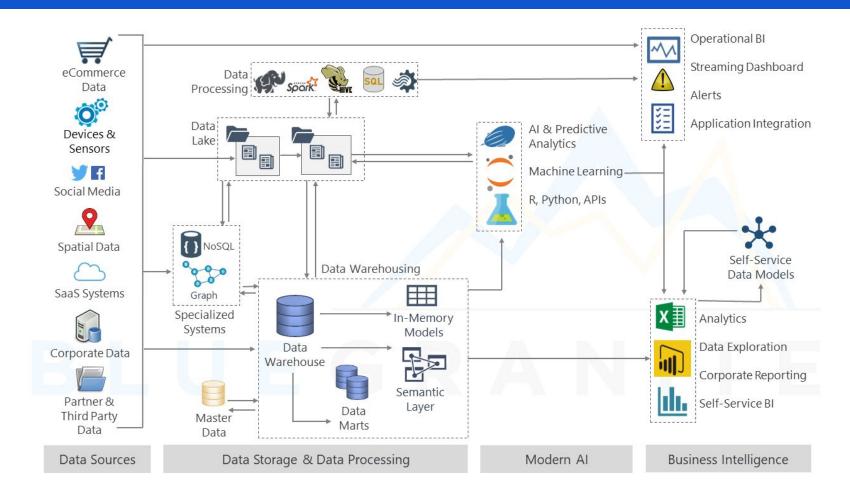








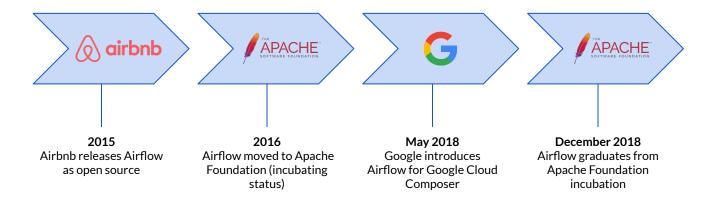
Modern Architecture



Apache Airflow

What Is Apache Airflow?

- Platform to programmatically author, schedule, and monitor workflows
- Developed by Airbnb in 2015, moved to Apache Foundation in 2016
- Used by 120+ companies, including Google, ING, Lyft, Paypal, Reddit, and more
- Useful for ETL, Machine Learning, Predictive, and General Pipelines



Airflow Principles & Features

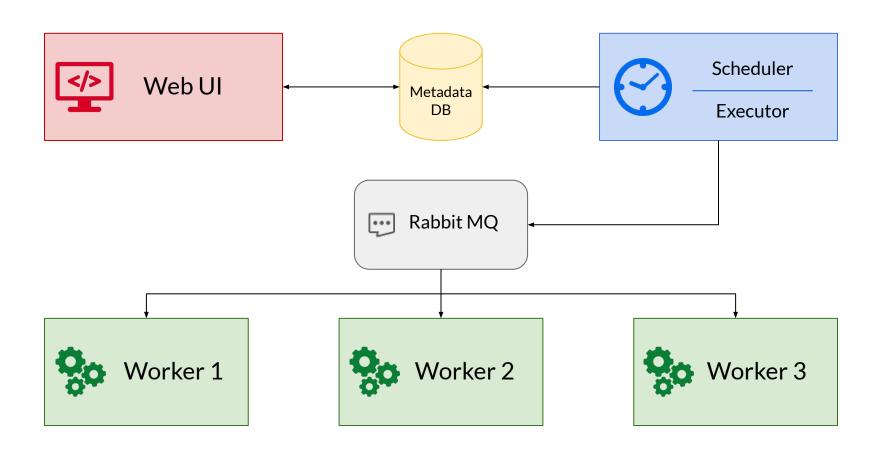
Principles

- Reproducible
 - Deterministic and idempotent
- Future proof
 - Backfilling, versioning
- Robust Against Changes
 - Easy to add, remove, modify DAGs
- Clarity
 - Transparency where data resides,
 what it means, and where it flows

Features

- Configuration as Code
- Usability (Web UI)
- Centralized Configuration
- Resource Pooling
- Extensibility
- Security (user, passwords)

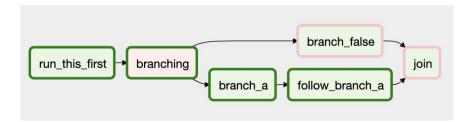
Airflow Architecture



Airflow Concepts

Airflow DAGs

"A DAG - or a Directed Acyclic Graph – is a collection of all the tasks you want to run, organized in a way that reflects their relationships and dependencies."



```
default_args = {
   'owner': 'Airflow',
   'depends_on_past': False,
    start_date': airflow.utils.dates.days_ago(2),
    email': ['airflow@example.com'],
   'email_on_failure': False,
   'email_on_retry': False,
   'retries': 1.
    retry_delay': timedelta(minutes=5),
     'queue': 'bash_queue',
     'pool': 'backfill',
     'priority_weight': 10,
      end_date': datetime(2016, 1, 1),
     'wait_for_downstream': False,
    'sla': timedelta(hours=2),
    'execution_timeout': timedelta(seconds=300),
     on failure callback': some function.
     'on_success_callback': some_other_function,
     'on_retry_callback': another_function,
     'sla_miss_callback': yet_another_function,
     'trigger_rule': 'all_success
dag = DAG('example',
          default_args=default_args,
          description='A simple tutorial DAG',
          schedule_interval=timedelta(days=1))
```

Airflow Operators

- BashOperator executes a bash command
- **PythonOperator** calls an arbitrary Python function
- EmailOperator sends an email
- **SimpleHttpOperator** sends an HTTP request
- SQL Operators executes a SQL command
- Sensor waits for a certain time, file, database row, S3 key, etc...

```
t1 = BashOperator(
  task_id='print_date',
  bash_command='date',
  dag=dag)
t2 = BashOperator(
  task_id='sleep',
  depends_on_past=False.
  bash_command='sleep 5',
  dag=dag)
templated_command = """
{% for i in range(5) %}
  echo "{{ ds }}"
  echo "{{ params.my_param }}"
t3 = BashOperator(
  task_id='templated',
  depends_on_past=False.
  bash_command=templated_command,
  params={'my_param': 'Parameter I passed in'},
  dag=dag)
t1 >> [t2, t3]
```

Airflow Operators

```
def print_context(ds, **context):
   """Print the Airflow context and ds variable from the context."""
   print(context)
   print(ds)
   return 'Whatever you return gets printed in the logs'
run_this = PythonOperator(
   task_id='print_the_context',
   provide_context=True,
   python_callable=print_context
   dag=dag)
def my_sleeping_function(random_base):
     "This is a function that will run within the DAG execution"""
   time.sleep(random_base)
# Generate 5 sleeping tasks, sleeping from 0.0 to 0.4 seconds respectively
for i in range(5):
   task = PythonOperator(
       task_id='sleep_for_' + str(i),
       python_callable=my_sleeping_function,
       op_kwarqs={'random_base': float(i) / 10},
       dag=dag)
   run this >> task
```

Airflow Operators

```
start = DummyOperator(task_id='start')
github = GithubToS30perator(task_id='github_to_s3'
                           github_conn_id=GITHUB_CONN_ID,
                           github_repo=GITHUB_REPO,
                           github_object=GITHUB_OBJECT,
                           s3_conn_id=S3_CONN_ID,
                           s3_bucket=S3_BUCKET,
                           s3_key=S3_KEY)
redshift = S3ToRedshiftOperator(task_id='s3_to_redshift',
                               s3_conn_id=S3_CONN_ID,
                               s3_bucket=S3_BUCKET,
                               s3_key=S3_KEY,
                               redshift_schema=REDSHIFT_SCHEMA,
                               redshift_conn_id=REDSHIFT_CONN_ID,
                               table='my_table')
start >> github >> redshift
```

Airflow Hooks

- Interfaces to external platforms and DBs
- Building block for operators
- Keep authentication code out of pipelines,
 centralized in the metadata DB

SQL (MSSQL, Oracle, Postgres)	AWS (S3, Lambda, Dynamo)
FS, FTP, SFTP, Samba, SSH	Azure (Cosmos, Data Lake, Fileshare)
HTTP, IMAP	Google Cloud (Big Table, ML, NLP)
Docker	Spark, Vertica
Redis, Mongo	Datadog, Salesforce, Slack

```
# Specify a connection
hook = SqliteHook('sqlite_connection_id')

# Get some records
records = hook.get_records('select * from accounts')
hook.run('delete from accounts')

# Bulk dump / load some data
hook.bulk_dump('accounts', '/tmp/accounts.csv')
hook.bulk_load('accounts', '/tmp/accounts.csv')

# Get a Pandas dataframe
df = hook.get_pandas_df('select * from users')

# Get a SqlAlchemy Engine
engine = hook.get_sqlalchemy_engine(...)
```

Airflow XComs

- XCom = Cross Communication
- Communication between tasks
- Can be "pushed" or "pulled" by all task instances

```
# Pushes an XCom without a specific target
def pusher1(**context):
  context['ti'].xcom_push(key='value 1',
                           value=value1)
# Pushes an XCom without a specific target,
# just by returning it
def pusher2():
  return 'value 2'
# Pull all previously pushed XComs
def puller(**context):
  ti = context['ti']
  # get value1
  value1 = ti.xcom_pull(key='value 1',
                         task_ids='pusher1')
  # get value2
  value2 = ti.xcom_pull(key='return_value',
                         task_ids='pusher2')
  # get both value1 and value2
  value1, value2 = ti.xcom_pull(
         key=None,
         task_ids=['pusher1', 'pusher2'])
```

Airflow Variables

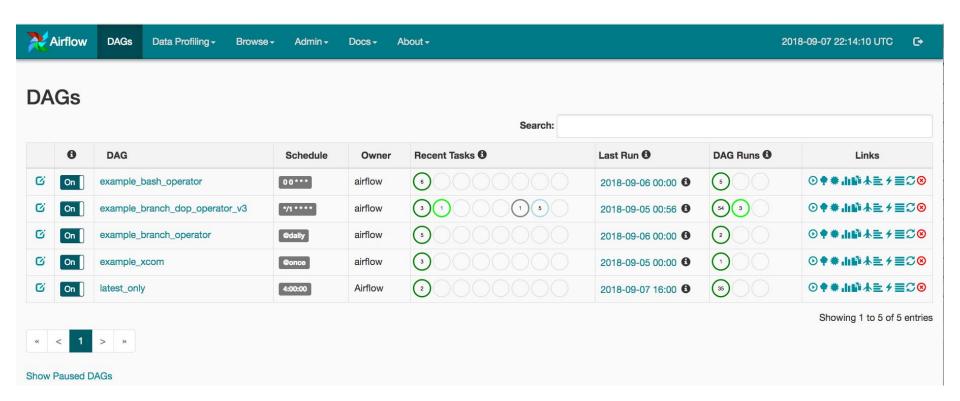
- Generic way to store and retrieve content as a simple key value store within Airflow
- Useful for configuration items that must be accessible and modifiable through a UI

```
# Variables can be in Python ...
from airflow.models import Variable
foo = Variable.get("foo")
bar = Variable.get("bar", deserialize_json=True)
baz = Variable.get("baz", default_var=None)

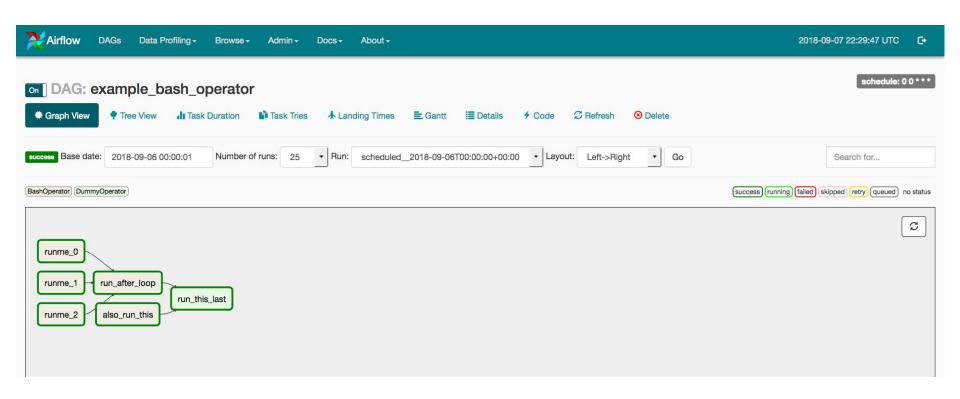
# ... or Jinja templates
echo {{ var.value.<variable_name> }}
echo {{ var.json.<variable_name> }}
```

UI Walkthrough

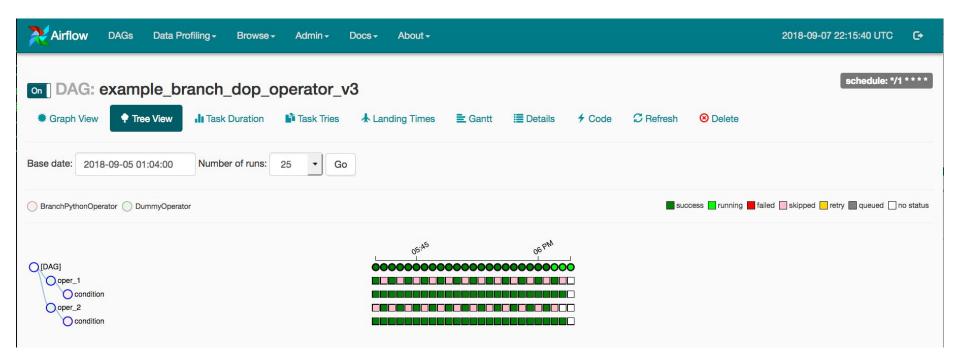
DAGS View



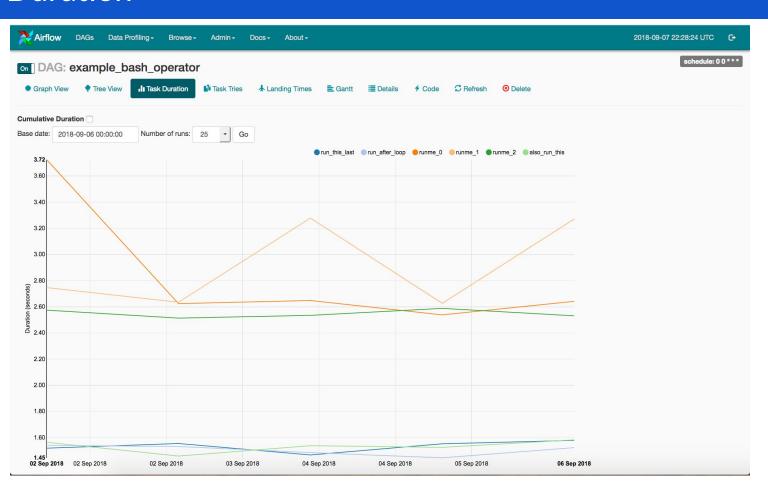
Graph View



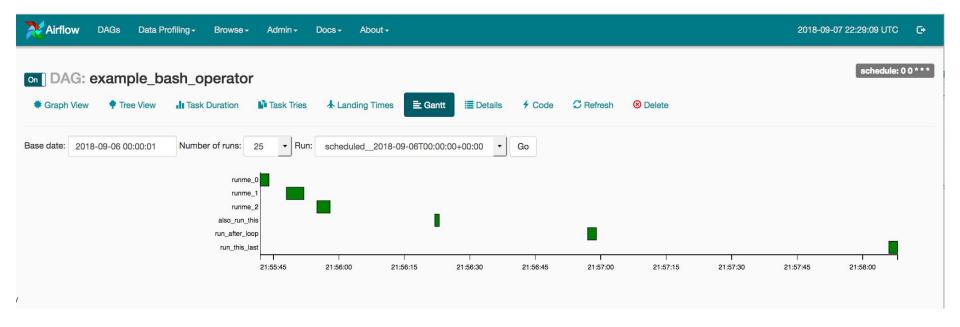
Tree View



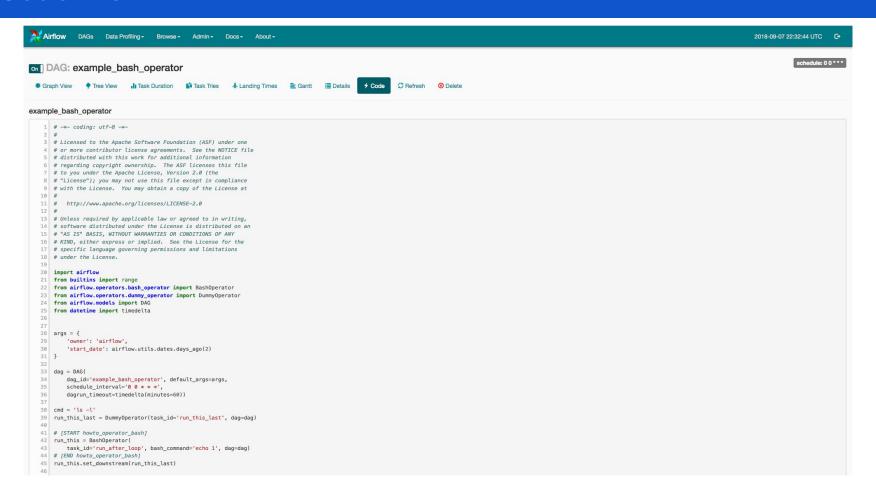
Task Duration



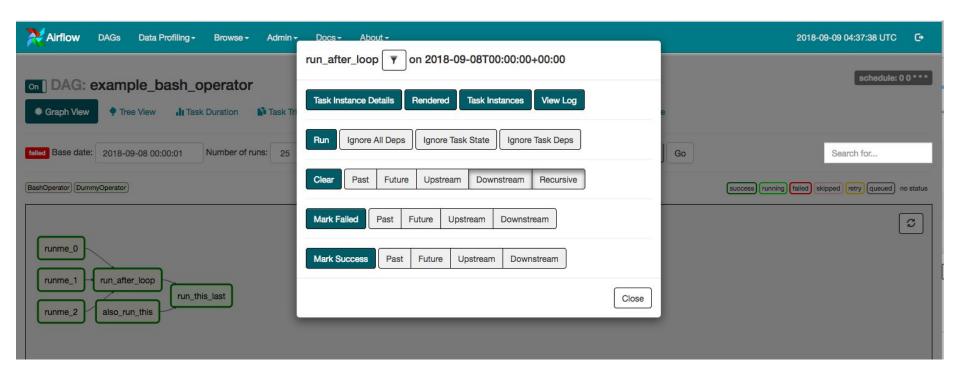
Gantt View



Code View



Task Instance View



Demos

Resources

Resources

Apache Airflow

https://airflow.apache.org/

Docker Apache Airflow

https://github.com/puckel/docker-airflow

Awesome Apache Airflow

https://github.com/jghoman/awesome-apache-airflow

Questions?